



Performance management 40 years on: a review. Some key decisions and consequences

Christopher Pollitt

To cite this article: Christopher Pollitt (2017): Performance management 40 years on: a review. Some key decisions and consequences, Public Money & Management, DOI: [10.1080/09540962.2017.1407129](https://doi.org/10.1080/09540962.2017.1407129)

To link to this article: <https://doi.org/10.1080/09540962.2017.1407129>



Published online: 22 Nov 2017.



Submit your article to this journal [↗](#)



Article views: 1



View related articles [↗](#)



View Crossmark data [↗](#)

Performance management 40 years on: a review. Some key decisions and consequences

Christopher Pollitt

This paper reviews the past four decades of experience of performance management (PM). Beginning with a brief history of its international spread and development, a short section on conceptual issues is followed by a strategic analysis of the research thus far. From this analysis is extracted a set of key decisions which must be made for any PM system. Each decision has significant consequences, and none are purely technical (though they may occasionally be represented as such). In many cases difficult trade-offs are inevitable.

Keywords: Indicators; international experience; key decisions; management; performance.

The story until now

In the UK some public services began to experiment with sets of performance indicators (PIs) in the late 1970s and early 1980s. One landmark was the appearance of the first national PI set for the National Health Service (NHS) in 1983 (Pollitt, 1985). Over the following decade PIs became ubiquitous—no self-respecting public service could be without them. Then under the Labour governments of 1997–2010 the performance indicator regime became more intensive, as more and more services faced penalties and adverse publicity if they failed to meet their targets. In local government, also, waves of PI systems became central (Martin *et al.*, 2013). The incoming Conservative/Liberal Democrat coalition of 2010 characterized this situation as one of undue centralization and red tape and promised ‘a bonfire of targets’. In practice, however, while buzzwords changed, most public managers did not notice much, if any, diminution of their responsibilities for detailed upward reporting of their organizational performance. At the time of writing, we can read almost daily newspaper accounts of the failures of UK hospitals, schools and prisons to reach their targets.

The US also developed a performance regime, starting, perhaps, a few years later than the UK. A landmark here was the Government Performance and Results Act 1992 (GPRA), although that was only one of a continuing stream of new performance information requirements that were laid upon federal agencies and managers (Radin, 1998).

Subsequently, at the top level, President G. W. Bush (2000–2008) made a Programme Assessment Rating Tool (PART) a central part of his management agenda—these were PIs by any other name. President Obama then tightened up the GPRA with the GPRA Modernization Act in 2010. At the sectoral level, New York Police Department’s Compstat and Baltimore’s Citystat became quite famous (Behn, 2014—see also the darkly humorous portrayal of Citystat in the excellent TV series *The Wire*). Indeed, a kind of ‘Compstatmania’ broke out across many city governments and federal agencies (Behn, 2014, pp. 1–11).

The performance movement spread to many other countries as well, including Australia, Canada, Denmark, Finland, France, The Netherlands, New Zealand, Norway and Sweden (Bouckaert and Halligan, 2008). Performance auditing by national and sub-national audit offices also flourished, especially in The Netherlands, Sweden and the UK (Pollitt *et al.*, 1999).

Furthermore, from the late 1990s, international indicator sets began to appear which compared broad features such as ‘effective governance’ or ‘red tape’ or corruption across many countries (for example Kaufmann *et al.*, 2009; Mungiu-Pippidi, 2015). By 2007, there were more than 400 such comparative indicator sets (IADB, 2007). Some academics began to look for ways to compare these different rating sets (Hood *et al.*, 2008), while others probed the politics behind international comparisons of good governance (Arndt, 2008).

Christopher Pollitt is Emeritus Professor, Leuven Institute of Public Governance, Belgium.

Indeed, academia responded to the fashion for PM with a growing tide of books and articles. These are by now far too many to list, but some key works by leading scholars have covered large theories of when PM works and doesn't (for example Boyne *et al.*, 2006; Moynihan, 2008; Hood and Dixon, 2010; Talbot, 2010), broad international comparisons (for example Bouckaert and Halligan, 2008; Hammerschmid *et al.*, 2013; Kuhlmann and Jäkel, 2013), analyses of specific policies in specific countries (Ingraham *et al.*, 2003; Perry *et al.*, 2009; Bevan and Wilson, 2013), and critiques of existing comparative governance indicators (for example Pollitt, 2011). There have also been numerous detailed case studies (for example Christensen *et al.*, 2006; Sundström, 2006; Bevan, 2009; Kelman and Friedman, 2009; Pollitt and Bouckaert, 2009; Pollitt *et al.*, 2010; Propper *et al.*, 2010). Despite many disappointments, there is evidence of PM having considerable impact across a range of jurisdictions (for example Kelman and Friedman, 2009; Bevan and Wilson, 2013; Behn, 2014). Several works have affirmed the central role played by PM in contemporary and future public sectors (Kettl and Kelman, 2007; Bouckaert and Halligan, 2008; Pollitt and Bouckaert, 2017). It is clear from these analyses that PM cannot be regarded solely as the lynchpin in the New Public Management (NPM) movement (as has occasionally been asserted). It has gone much wider than that and has been used, in a variety of ways, by a wide range of approaches to reform and improvement.

A brief conceptual interlude

This is not the place for a full elaboration of the several conceptual schemes and distinctions which some academics have developed around PM (for example Bouckaert and Halligan, 2008, pp. 15–38; Talbot, 2010, pp. 17–18). Instead there will be offered no more than a quick introduction to the concept of performance itself, together with brief working definitions of our two main terms (PIs and PM), and observations thereon.

'Performance' is not a precise concept, nor a scientific term. In normal usage it can refer to the process ('He will get more teaching hours than at his previous school') or the outputs ('She got grade As for all her A-levels'), or the final outcomes ('Without that change of school she would never have become the top lawyer she is today'). In the public sector context performance is obviously multi-dimensional—a 'good performance' could

refer to cost, efficiency, effectiveness, freedom from corruption, public access, quality of service or a host of other dimensions. Thus it is always worth enquiring which definition of performance is actually operationally in use, and who has the power to craft and to impose that definition. Equally, it is easy for politicians and managers to refer to 'improved' (or 'declining') performance in a vague, general way, skating over the fact that certain aspects of performance may be improving while others are simultaneously declining. Beyond that, there is usually an unwillingness, on behalf of politicians especially, to acknowledge the likelihood that a planned improvement on one performance dimension will probably lead to the deterioration of another (as when an economy drive lengthens waiting times or reduces overall effectiveness).

It should be noted that this paper concentrates on *organizational* performance—the measured achievements of departments, agencies, units and so on. The performance of *individuals* is clearly also important, and is dealt with in a separate but connected literature, but is not pursued here. The multi-dimensionality of organizational performance is present even in the simplest public service, but it is amplified hugely in those complex public services—such as healthcare, education or the police—which make up the backbone of most western public sectors. A hospital, school or police force simultaneously does so many *different* things for their publics.

In this context, a PI is a measure of some aspect or dimension of performance—such as the average waiting time in a hospital's accident and emergency (A&E) department. It may be cardinal, ordinal or nominal. However, there is not much point to a PI or set of PIs unless it is part of a wider system of performance management (PM). (Despite this, the literature does record instances where PI results are quite detached from management processes and are usually just filed and forgotten—or kicked around as political and media footballs.) PM simply means that there is a deliberate, conscious attempt on the part of organizational management to steer the organization in directions that will improve its PI scores. Behn (2014) conceives this as a leadership strategy. Such steering may involve exhortations and information (feedback), recruitment and training, the reward or sanctioning of members of staff, the re-allocation of resources or a range of other instruments. Above all, it needs

to involve persistent *discussion* of measures, targets, results and tactics.

Note also that an indicator may be simple or composite. A simple indicator might be the average time taken to process an application and issue a license. A composite indicator might be a weighted average of eight key PIs of a surgical unit in a hospital. Composite indicators have considerable attractions in enabling PI users to 'see the big picture'; they can be used to sum up a variety of different aspects of the performance of a complex service. Composites may also have several drawbacks (Jacobs *et al.*, 2006). One is that they may smuggle in disguised political choices via the apparently technical weighting system. They can also become *so* complex that it is very difficult to understand what the composite figure really means (Pollitt, 2011).

Finally, two further problems should be acknowledged. The first is that the same PI set—or individual PI—is frequently assessed differently by different stakeholders. The conditions that need to be present before different stakeholders arrive at the same or similar judgements are quite demanding (Boyne *et al.*, 2006, pp. 18–19). The second is that the conditions necessary for citizens to trust performance information are also challenging—and often are not met (Pollitt and Chambers, 2013).

Decisions and consequences

There are a multitude of different approaches to the task of extracting the essentials from the vast mound of studies of PM and PIs which already exists (and which is still being enthusiastically added to). The method here is to identify a relatively small number of key decisions that inevitably *have* to be made in the design of *any* PM system, and then to examine the consequences of making each of these one way or another. In most cases these decisions entail trading off one or more desiderata against others. By this method, therefore, the underlying structure of PM systems can be delineated and explored. Its application swiftly reveals that any idea that PI systems can be made fully comprehensive, totally objective (de-politicized) and somehow automatic is and always has been a fantasy. PIs can be enormously helpful and illuminating, but they cannot be purely technical. Technological advance can and has made the collection, processing and dissemination of performance data immeasurably easier than it was 40 years ago, but it cannot resolve the judgemental challenges embodied in the issues we are about

examine. Design decisions are political and managerial decisions as well as technical ones, and they certainly have political and managerial consequences. As one commentator put it:

Rankings of these performances will always reflect a certain dominant definition of performance. Ranking schemes as a result both reflect and create performance (Van de Walle, 2008, p. 256).

Decision: Who is it for?

There are many potential answers to this. Politicians? Top managers? Operational managers? Service users? The general public? The easy, diplomatic but usually unrealistic, answer is to say 'For all of these'. This is usually unrealistic because, as many studies have shown, PIs designed to meet the priorities of one stakeholder are unlikely to meet the priorities of all the others (Behn, 2003). Of course different sets of PIs may be included—this set for the service users, that one for the operational staff, and so on—but that route quickly generates a large and unwieldy system which may well be hard to manage. There are limits to what a single, manageable PM system can do. It usually cannot be fine-tuned to meet the needs of all the various stakeholders, so some tough decisions have to be made about who it is really for. Making it useful to more than one stakeholder may be perfectly feasible, but making it equally valuable for all is a much taller order.

Closely allied to this question is that of ownership. Many studies show that PM systems work best if the operational staff 'own' them—recognize their legitimacy and pay them attention. But if the system was imposed on them, perhaps by top management in order better to control them, or by politicians in the name of greater 'accountability', then legitimacy with operational staff may be hard to secure. PM is more difficult to do if the operational staff (who are often the recorders as well as the expected users of the data) are in a state of recalcitrance.

Certain stakeholders seem particularly difficult to reach. Politicians and the public emerge from several studies as hard to engage. Or, rather, it may be easy to reach them but it is very hard to hold their attention and gain their understanding (Pollitt, 2006a; Johnston and Talbot, 2008; Pollitt and Chambers, 2013). In both cases the information has to be simple and right in front of them precisely when their need for it arises (for example performance data on local schools when parents are choosing

where to send their children). Whilst there have been some promising initiatives with both groups (for example Ho and Coates, 2004) widespread engagement has thus far proved elusive.

Decision: What should be measured?

This fundamental decision is closely influenced by the previous question. But in all cases, whoever the 'owners' are, there will always be some selection—some prioritizing—to be made. Usually, what gets measured gets attention, so what does not get measured may become relatively neglected (Smith, 1995; Bevan and Hood, 2006). The notorious case of the Mid Staffordshire NHS Foundation Trust hospital, where many A&E patients suffered and died prematurely, arose at least partly because the hospital management focused mainly on PIs for throughput, while quality of clinical care issues were backgrounded (Healthcare Commission, 2009).

Actually, in the early days, and even now, many PIs are not even about outputs. They are measures of internal organizational processes, such as hospital length of stay. True output and outcome measures are often in short supply. In an analysis of 519 studies of NPM reforms from 26 EU member states, it was found that only 9% contained information about outcomes, with 27% containing information about outputs and/or outcomes (Pollitt and Dan, 2013).

To return to the output/outcome balance, the main reason why it has proved controversial is that, while outcomes are clearly the gold standard (what is actually *achieved*), they are frequently difficult to measure and/or difficult to attribute with any confidence to the outputs of the public service concerned (Behn, 2014, p. 198). Thus, for example, one might wish to measure the performance of the government's job assistance programmes by the levels of employment/unemployment in the economy, but changes in these levels are mainly influenced by factors far beyond the control of state job agencies. Not unreasonably, doctors and nurses do not wish to be held responsible for the health consequences of changes in the nation's eating or drinking habits, or for bad housing. So service delivery agencies cannot necessarily be held accountable for outcomes (changes in health statuses, employment levels, crime rates etc.), whereas they can be held responsible for their own outputs (licences issued, lessons delivered, health treatments provided, arrests and convictions etc.). The implication of this state of affairs is that PM

systems based *exclusively* on either outcomes or outputs are likely to be unsatisfactory. Some mixture is required because we want to know both about the quality of what service organizations are providing and about what ultimate effects (if any) are being registered in the outside world. It is clear, however, that PM systems consisting mainly of process measures will be disappointing for most external stakeholders, including politicians and public.

Longitudinal studies have shown that PI sets tend to become more complex over time. Individual indicators are refined, but new indicators are also added, as voices call for additional dimensions of performance to be captured rather than neglected. For example, in the case of the NHS, indicators of clinical outcomes and indicators of patient satisfaction were gradually added to the national set over the years (Pollitt *et al.*, 2010). The numbers of PIs rise and even if there are occasional attempts to cut back to simpler, core sets of measures these reductions seldom go as far back as square one. In short, the original decision as to what should be measured seldom remains set in concrete over time.

Decision: Who will measure?

A related question is that of who carries out the measurement? If it is the service providers themselves, there are obvious dangers of self-serving distortions to the data. But if independents are to be brought in to do the measurement that may be cumbersome, expensive and corrosive of the operational staff's ownership and trust. The literature contains a number of cases in numbers being gamed or even falsified (Audit Commission, 2003; Ofqual, 2012; Pollitt, 2013). External validation may be at least a partial answer to this but, again, it can be cumbersome. On several occasions, the UK National Audit Office pressed the government to have its major PI sets independently validated, but this proposal was never accepted. This looks very much like a trade-off.

Connected questions are how data will be presented, when and to whom? Again, there are bound to be issues of self-interest, transparency and trust that go far beyond the technical. Elaborate tables of composite indicators, however methodologically sound, will fail to hold the attention of many politicians, citizens and even operational staff. Different formats are required for different audiences. Data that is not timely will also fail to impact: few want to read about how an agency performed a year or more ago—they want to

know how it is performing now. More generally, in the US and the UK particularly, providing the public with trusted data has been problematic, because in the minds of many citizens little if anything from the government can be trusted. Independent measurement, clearly removed from government, may be at least a partial answer to this, which returns us to the trade-off with the trust of operational staff.

Decision: What criteria will be applied to judge outstanding, acceptable and poor performances?

Equally important with the raw numbers yielded by the measurements are the criteria that are applied to them. For example, a given performance can be compared against (a) a professional standard; or (b) an average or mean; or (c) a threshold set by the public authorities ('No patient shall wait more than 48 hours to see their primary care physician'); or (d) the best performance of any organizational unit in the regional or national (or international) set under consideration; or (e) some combination of the above. The same score may look creditable when measured against one criterion, yet it may appear poor when set against another.

Much follows from the earlier decision about who the system is primarily for. The chosen criteria should presumably be relevant and acceptable to that stakeholder/those stakeholders. Unfortunately, that may mean that they may not be well-regarded by other stakeholders. Neurosurgeons and their patients, university professors and their students, police and public—all these pairs and more may have differing views of what criteria are appropriate, or may simply not know enough to set a sensible criterion. Taking these decisions is not so much trading off as negotiating around what will be meaningful, comprehensible and useful to the various parties. Surprisingly little seems to have been written about this.

Criteria may change over time, even if the PIs remain the same. If a threshold of acceptability is initially set at $x\%$ and after a few years most operational units have achieved at least x , it makes perfect sense to raise the criterion to $x + n$. Note, however, that the service needs to be well understood before such increases are implemented—otherwise those in authority may find themselves asking for the impossible. Better organization and new technology may enable hospitals to reduce the average length of stay for a particular

surgical procedure from six days to three. But that does not mean that the target can then be reduced from three to zero. There is an underlying need to understand the specific process leading from actions to output and then connecting the output to identifiable outcomes (Pawson, 2013; Behn, 2014, pp. 145–171). This often requires the toughest analytic thinking.

Decision: What incentives/penalties (if any) will be attached to the PIs?

Arguments about incentives have raged since the earliest days of PIs. One position is that no incentives should be attached, so that professional service providers such as doctors and teachers can deliberate among themselves on the meaning of the scores and can then select the most appropriate responses, without fear of sanction. At the other end of the scale lies the argument that, unless real and visible penalties and rewards are attached to poor or outstanding performance, service providers will pay little attention to them. A revealing book by the former head of the (UK) Prime Minister's Delivery Unit under Blair exhibited sentiments much closer to the latter argument than the former. Professionals—doctors, teachers, the police and others—were not to be trusted to deliver improvement unless real pressure was applied (Barber, 2007). Radin's work on the performance movement in the US has a whole chapter entitled 'Demeaning professionals' (Radin, 2006). Comparative work has shown that very similar PIs may be applied differently in some cultures than others—more 'softly', say, in the Nordic countries than in the UK (Pollitt, 2006b).

Under the Blair government the UK furnished an interesting natural experiment for testing these arguments. In hospital care England linked PIs to a system of strong incentives and penalties. In Scotland and Wales, this linkage was nowhere near as strong—drawing lessons from the PIs was left more to the professionals. Detailed studies indicate that performance improvement was much greater in England than in Scotland or Wales (Connolly *et al.*, 2009; Propper *et al.*, 2010). However, there are always questions of what other aspects of performance may be being neglected in order to hit headline targets (Bevan and Hood, 2006), and occasionally there are clear cases where strong incentives have resulted in highly undesirable behaviours (Francis, 2010; Pollitt, 2013).

Which again brings us to a difficult trade-off. If there are no incentives then behaviour

may not change much, and performance will remain at its previous level. If there are very powerful incentives then perverse behaviours tend to appear as staff try to game the system in order to avoid punishment (Bevan and Hood, 2006; Bevan, 2009; Pollitt, 2013).

One study of the development of the English and Dutch systems of healthcare PIs over the period 1982–2007 suggested that there might be an inherent ‘logic of escalation’ present in PM (Pollitt *et al.*, 2010). Two elements in this were a shift from formative to incentive-linked summative measures and a parallel shift from intra-professional PIs to PI sets which were deliberately disseminated to the public and the media (sometimes called ‘naming and shaming’—Bevan and Wilson, 2013). As soon as a scandal comes along, it is hard for politicians and top managers to resist the call to make PI sets open and to link them to penalties for the poorest performers.

Decision: How often should PIs be changed?

This may not at first sight seem a particularly important question, but research across different countries and sectors indicates that it often is. The reason for this is, once more, an underlying trade-off. If PIs (or their criteria) remain fixed for a long time their effect weakens. Staff learn tricks to work round them. Organizations which have achieved a satisfactory performance may have little incentive to improve further (the ‘threshold effect’—see Pollitt, 2013). The distribution of performances across the relevant population of organizations will narrow so that there are no longer such big differences between the best and the worst, and the PI therefore becomes less discriminating. The gap between the underlying ‘real’ performance and the apparent scores will widen—this tendency has been dubbed the ‘performance paradox’ (Meyer and Gupta, 1994; Van Thiel and Leeuw, 2002). On the other hand, if PIs are changed very frequently, operational staff may become confused or despairing—it is hard to make medium or longer term plans if every year your targets change. So, in broad terms, it seems sensible to aim for some ‘churn’ of PIs, but nothing too hectic. This is easier said than done. There is no formula for calculating the optimal rate of change and, in any case, this is likely to vary considerable according to the resource climate and the characteristics of a particular service (some services are inherently more

short/long term in their operations than others, and may require greater continuity).

Conclusions

The now-voluminous PM literature yields a reasonably consistent picture. Certain key issues recur over time and across different sectors and countries. In the above analysis, these have been rendered into the form of questions about the design of PM systems. Of course these are not the only issues that arise with PM, but they are fundamental, for the reasons given above. None of these decisions has a single, correct answer. Most of them are questions of where to strike a balance between different objectives which are in tension one with another. Some interact—for example, decisions about who the system is primarily for will have implications for what incentives (and penalties) seem appropriate. Neither can it be assumed that the optimal answers for one service or set of circumstances will fit another organization, time or place. Contextual influences are frequently important, with similar PIs being used in quite different ways in different locations, systems and, even more, cultures (Pollitt, 2006b; De Kool, 2008; Pollitt and Dan, 2013; Behn, 2014, pp. xvi and 31–35).

PM systems are therefore definitely not ‘fire and forget’. They require active maintenance and continuing reconsideration. They require tuning to local circumstances. When the six key decisions are taken intelligently they can bring very significant benefits. When these decisions are taken badly, or even ignored, PM becomes less useful—sometimes grindingly bureaucratic, sometimes merely decorative and sometimes actively harmful. PM systems are also attended by a whole suite of technical questions—how to measure this or that activity accurately, how new data collection systems should be organized, and so on, but if the fundamental decisions are mistaken no amount of technical perfection can compensate. An approximate measure of an important aspect of performance is far more valuable than a precise measure of something trivial.

Finally, it should be borne in mind that PM systems have experienced frequent failures and disappointments, alongside their many solid achievements. The appropriate stance of the promoters and designers of these systems should perhaps be more one of humility than of hubris. They should bear in mind in their own work one of the

foundational lessons that they often try to teach to others. As the poet, Leigh Hunt, put it in 1832:

The pretension is nothing; the performance everything.

References

- Arndt, C. (2008), The politics of governance ratings. *International Public Management Journal*, 11, 3, pp. 275–297.
- Audit Commission (2003), *Waiting List Accuracy* (The Stationery Office).
- Barber, M. (2007), *Instruction to Deliver: Tony Blair, Public Services and the Challenge of Achieving Targets* (Politicos).
- Behn, R. (2003), Why measure performance? Different purposes require different measures. *Public Administration Review*, 63, 5, pp. 586–606.
- Behn, R. (2014), *The Performance Stat Potential: A Leadership Strategy for Producing Results* (Brookings Institution Press).
- Bevan, G. (2006), Setting targets for healthcare performance: lessons from a case study of the English NHS. *National Institute Economic Review*, 197, 1, pp. 67–79.
- Bevan, G. (2009), Hitting and missing targets by ambulance services for emergency calls: effects of different systems of performance measurement in the UK. *Journal of the Royal Statistical Society*, 172, 1, pp. 161–190.
- Bevan, G. and Hood, C. (2006), What's measured is what matters: targets and gaming in the English public healthcare system. *Public Administration*, 84, 3, pp. 517–538.
- Bevan, G. and Wilson, D. (2013), Does 'naming and shaming' work for schools and hospitals? Lessons from natural experiments following devolution in England and Wales. *Public Money & Management*, 33, 4, pp. 245–252.
- Bouckaert, G. and Halligan, J. (2008), *Managing Performance: International Comparisons* (Routledge).
- Boyne, G., Meier, K., O'Toole Jnr, L. and Walker, R. (Eds), *Public Service Performance: Perspectives on Measurement and Management* (Cambridge University Press).
- Christensen, T., Laegreid, P. and Stigen, I. (2006), Performance management and public sector reforms: the Norwegian hospital reform. *International Public Management Journal*, 9, 2, pp. 1–27.
- Connolly, S., Bevan, G. and Mays, N. (2009), *Funding and Performance of Healthcare Systems of the Four Countries of the UK Before and After Devolution* (The Nuffield Trust).
- De Kool, D. (2008), Rational, political and cultural uses of performance monitors: the case of Dutch urban policy. In Van Dooren, W. and Van de Walle, S. (Eds), *Performance Information in the Public Sector: How it is Used* (Palgrave Macmillan), pp. 174–191.
- Francis, R. (2010), *Independent Inquiry into the Care Provided by Mid Staffordshire NHS Foundation Trust, January 2005–March 2009*, HC375-1 (The Stationery Office).
- Hammerschmid, G., Van de Walle, S. and Stimac, V. (2013), Internal and external use of performance information in public organizations: results from an international survey. *Public Money & Management*, 33, 4, pp. 261–268.
- Healthcare Commission (2009), *Investigation into Mid Staffordshire NHS Foundation Trust*.
- Ho, A. and Coates, P. (2004), Citizen-initiated performance assessment: the initial Iowa experience. *Public Performance and Management Review*, 27, 3, pp. 29–50.
- Hood, C., Dixon, R. and Beeston, C. (2008), Rating the rankings: assessing international rankings of public service performance. *International Public Management Journal*, 11, 3, pp. 298–328.
- Hood, C. and Dixon, R. (2010), The political payoff from performance target system: no-brainer or no-gainer? *Journal of Public Administration Research and Theory*, 20, pp. i281–i298.
- Hunt, L. (1832), *The Poetical Works of Leigh Hunt* (Ward Lock & Co.).
- Inter-American Development Bank (IADB), (2007), *Datagob: Governance Indicator Database* (<http://www.iadb.org/datagob/>).
- Ingraham, P., Joyce, P. and Donahue, A. (2003), *Government Performance. Why Management Matters* (John Hopkins University Press).
- Jacobs, R., Goddard, M. and Smith, P. (2006), *Public Services: Are Composite Measures a Robust Reflection of Performance in the Public Sector?* (Centre for Health Economics, University of York), Research paper 16.
- Johnston, C. and Talbot, C. (2008), UK parliamentary scrutiny of public service agreements: a challenge too far? In Van Dooren, W. and Van de Walle, S. (Eds), (2008), *Performance Information in the Public Sector: How it is Used* (Palgrave Macmillan), pp. 157–173.
- Kaufmann, D., Kraay, A. and Mastruzzi, M. (2009), *Governance Matters VIII: Aggregate and Individual Governance Indicators for 1996–2008* (World Bank), Policy research working paper 4978.
- Kelman, S. and Friedman, J. (2009), Performance improvement and performance dysfunction: an empirical examination of the distortionary impacts of the emergency room wait-time target

IMPACT

The paper reviews the now vast international literature on performance management (PM) and performance indicators (PIs). It picks out six key decisions that are inevitably faced by anyone designing a PM system. Each decision is examined in the light of the evidence. Most of them appear to entail some kind of trade-off. It is imperative that practitioners should be aware of these unavoidable decisions and their likely consequences. It is also vital that PM designers should take account of local circumstances and contextual factors.

in the English National Health Service. *Journal of Public Administration Research and Theory*, 19, 4, pp. 917–946.

Kettl, D. and Kelman, S. (2007), *Reflections on 21st Century Government Management* (IBM Center for the Business of Government).

Kuhlmann, S. and Jäkel, T. (2013), Competing, collaborating or controlling? Comparing benchmarking in European local government. *Public Money & Management*, 33, 4, pp. 269–276.

Martin, S., Downe, J., Grace, C. and Nutley, S. (2013), New development: All change? Performance assessment regimes in UK local government. *Public Money & Management*, 33, 4, pp. 277–280.

Meyer, J. and Gupta, V. (1994), The performance paradox. *Research in Organizational Behavior*, 16, pp. 309–369.

Moynihan, D. (2008), *The Dynamics of Performance Management* (Georgetown University Press).

Mungiu-Pippidi, A. (2015), *The Quest for Good Governance: How Societies Develop Control of Corruption* (Cambridge University Press).

Ofqual (2012), *GCSE English 2012*, Ofqual Report 12/5225.

Pawson, R. (2013), *The Science of Evaluation: A Realist Manifesto* (Sage).

Pollitt, C. (1985), Performance indicators: can practice be made perfect? *Health and Social Services Journal* (6 June), pp. 706–707.

Pollitt, C. (2006a), Performance information for democracy: the missing link? *Evaluation*, 12, 1, pp. 38–55.

Pollitt, C. (2006b), Performance management in practice: a comparative study of executive agencies. *Journal of Public Administration Research and Theory*, 16, 1, pp. 25–44.

Pollitt, C. (2011), ‘Moderation in all things’: international comparisons of governance quality. *Financial Accountability & Management*, 27, 4, pp. 437–457.

Pollitt, C. (2013), The logics of performance

management. *Evaluation*, 19, 4, pp. 346–363.

Pollitt, C., Girre, X., Lonsdale, J., Mul, R., Summa, H. and Waerness, M. (1999), *Performance or Compliance? Performance Audit and Public Management in Five Countries* (Oxford University Press).

Pollitt, C. and Bouckaert, G. (2009), *Continuity and Change in Public Policy and Management* (Edward Elgar).

Pollitt, C., Harrison, S., Dowswell, G., Jerak-Zuiderent, J. and Bal, R. (2010), Performance regimes in healthcare: critical junctures and the logic of escalation in England and the Netherlands. *Evaluation*, 16, 1, pp. 13–29.

Pollitt, C. and Chambers, N. (2013), Evidence based trust: a contradiction in terms? In Llewellyn, S., Brooks, S. and Mahon, A. (Eds), *Trust and Confidence in Government and the Public Services* (Routledge), pp. 36–45.

Pollitt, C. and Dan, S. (2013), Searching for impacts in performance-oriented management reform: a review of the European literature. *Public Performance and Management Review*, 37, 1, pp. 7–32.

Pollitt, C. and Bouckaert, G. (2017), *Public Management Reform: A Comparative Analysis—Into the Age of Austerity*, 4th edn (Oxford University Press).

Propper, C., Sutton, M., Whitnall, C. and Windmeijer, F. (2010), Incentives and targets in hospital care: a natural experiment. *Journal of Public Economics*, 94, 3, pp. 301–335.

Radin, B. (1998), The Government Performance and Results Act (GPRA): hydra-headed monster or flexible management tool? *Public Administration Review*, 58, 4, pp. 307–316.

Radin, B. (2006), *Challenging the Performance Movement: Accountability, Complexity, and Democratic Values* (Georgetown University Press).

Smith, P. (1995), On the unintended consequences of publishing performance data in the public sector. *International Journal of Public Administration*, 18, pp. 277–310.

Sundström, G. (2006), Management by results: its origin and development in the case of the Swedish state. *International Public Management Journal*, 9, 4, pp. 399–427.

Talbot, C. (2010), *Theories of Performance* (Oxford University Press).

Van Thiel, S. and Leeuw, F. (2002), The performance paradox in the public sector. *Public Performance and Management Review*, 25, 3, pp. 267–281.

Van de Walle, S. (2008), What services are public? What aspects are to be ranked? The case of ‘services of general interest’. *International Public Management Journal*, 11, 3, pp. 256–274.